

# A Unified Approach to Pose Estimation in Elephants and Other Quadrupeds using Noisy Labels

Karen Panetta<sup>1</sup>, Obafemi Jinadu<sup>1\*</sup>, Jamie Heller<sup>1</sup>, Srijith Rajeev<sup>1</sup>, Kali Pereira<sup>2</sup>, Sos Agaian<sup>3</sup>

<sup>1</sup> School of Engineering, Tufts University, Medford, Massachusetts, 02155, United States

<sup>2</sup> Cummings School of Veterinary Medicine, Tufts University, North Grafton, Massachusetts, 01536, United States

<sup>3</sup> Department of Computer Science, College of Staten Island, The City University of New York, New York, 10314, United States

Correspondence and requests for materials should be addressed to O.J. (email: [obafemi.jinadu@tufts.edu](mailto:obafemi.jinadu@tufts.edu))

## Summary

Pose estimation predicts anatomical landmarks in humans and animals from monocular images or videos. Animal Pose Estimation is crucial for monitoring locomotion, behavior, and activity recognition, playing a key role in wildlife conservation. Single species pose estimation studies capture features unique to the species but generalize sub-optimally, while multi-species studies provide broader generalization by assuming fixed keypoints for all quadrupeds, this oversimplification fails to capture unique anatomical traits in animals such as elephants. To harness the strengths of single-species and multi-species pose estimation, we present QuadPose, a framework that standardizes skeletal structures across datasets and improves generalizability through consistency-dependent pseudo-labelling. Additionally, JumboPose, a manually annotated dataset of 2,078 African elephant images with 33 keypoints tailored to their unique morphology is introduced. Extensive evaluations demonstrate the effectiveness of QuadPose for animal pose estimation. This work establishes a foundation for standardized, cross-species pose estimation, advancing applications in wildlife conservation, and veterinary research.

## Introduction

The field of computer vision has tremendously benefited from advances in machine learning. State-of-the-art methods for object detection<sup>1,2,3,4</sup>, recognition<sup>5</sup>, segmentation<sup>6,7,8,9,10</sup> and pose estimation<sup>11,12,13,14,15,16,17</sup> adopt deep neural network frameworks, such as convolutional neural networks<sup>18,5</sup> and transformer-based architectures<sup>19,20</sup> for feature extraction. Pose estimation involves identifying and localizing anatomical keypoints such as elbows, wrists<sup>11</sup> within an image or video. Pose estimation methods are generally categorized into top-down<sup>21</sup> and bottom-up<sup>22,23</sup> approaches. The top-down approach involves using a detector model<sup>1,2,24,25,26</sup> to isolate regions of interest in the image by obtaining a set of coordinate locations for each detected object and subsequently performs single-person pose estimation. The bottom-up approach does not require a separate detector; instead, it directly predicts keypoints all at once, followed by an association step that groups them into full poses for each individual<sup>27,28,29</sup>.

Practical applications that require pose estimation as a critical step include behavior understanding, human-object interaction, and activity recognition. The availability of large-scale datasets, such as COCO<sup>30</sup>, OCHuman<sup>31</sup> and MPII<sup>32</sup>, has been a driving force behind advancements in human pose estimation (HPE). Despite advances in HPE, progress in animal pose estimation (APE) remains limited. A major challenge is the lack of large-scale, labeled datasets that comprehensively represent diverse animal species. Without standardized datasets and robust models, APE systems struggle to accurately track animal movement, monitor health conditions, and support conservation efforts in real-world settings. To address this, several datasets have been created for specific quadrupeds, with annotations tailored to their anatomical structure. These

47 include datasets for tigers<sup>33</sup>, horses<sup>34</sup>, macaque<sup>35</sup>, and zebra<sup>36</sup>. Multi-species datasets, such as  
48 Animal-Pose<sup>37</sup> and AP-10K<sup>38</sup> with five and fifty-four animal species respectively, aim to capture  
49 broader quadruped similarities. These datasets have gained traction in wildlife conservation  
50 research. However, annotation styles and skeletal structures across these APE datasets vary. These  
51 differences include the number and placement of keypoints, inconsistent naming conventions, and  
52 varying anatomical definitions. These discrepancies reduce knowledge transferability in large-  
53 scale cross dataset learning. Despite these inconsistencies, the shared anatomical similarities among  
54 quadrupeds present an opportunity to standardize and consolidate images and labels across most  
55 of these datasets into a unified framework, leveraging common features across species. These  
56 anatomical similarities can be extended to out-of-domain classes to improve the model’s  
57 generalizability to unseen animal species.

58 Another challenge worth addressing is the failure of existing datatsets<sup>37,38,35,33,36,34</sup> to  
59 capture critical keypoints in certain quadrupeds, such as elephants, which have unique features not  
60 shared by other species, including trunks, large ears, and tusks. This underscores the need for a  
61 framework capable of reliably estimating poses for quadruped species with distinct morphologies,  
62 such as elephants, while also generalizing effectively to other quadrupeds.

63 This work, introduces QuadPose, a unified framework for animal pose estimation that  
64 enables accurate, species-specific predictions for elephants while improving generalization across  
65 diverse quadrupeds. To address the lack of datasets capturing the unique morphology of elephants,  
66 we propose JumboPose, a large-scale dataset dedicated to elephant pose estimation, featuring  
67 2,078 manually annotated images with 33 anatomically relevant keypoints.

68 QuadPose formulates APE as a multi-task learning problem, standardizing pose estimation  
69 into two data representations: one tailored specifically for African elephants and another  
70 encompassing all other quadrupeds, implemented as a dual-head prediction network. Our  
71 framework is developed based on top-down state-of-the-art architectures, including HRNet<sup>18</sup> with  
72 polarized self-attention<sup>39</sup>, ViTPose<sup>40</sup>, and TransPose<sup>41</sup>, and incorporates a binary classifier that  
73 dynamically routes input data to the appropriate prediction head. Additionally, we introduce a  
74 pseudo-labeling strategy that leverages shared anatomical features to enhance generalization to  
75 unseen animal species. Extensive evaluations demonstrate that QuadPose achieves state-of-the-art  
76 performance, with mAP scores of 81.5, 85.7, and 94.3 on Animal-Pose, AP-10K, and JumboPose,  
77 respectively. By standardizing skeletal structures and leveraging multi-task learning, QuadPose  
78 not only enhances species-specific accuracy but also improves cross-species generalization. This  
79 framework establishes a new benchmark for scalable and robust animal pose estimation, paving  
80 the way for broader applications in wildlife conservation, behavioral analysis, and veterinary  
81 science. In particular, these capabilities hold significant potential for monitoring free-roaming  
82 wildlife populations, enabling automated censusing by age-sex class and facilitating the remote  
83 detection of sick or injured individuals. Such advancements are crucial for improving conservation  
84 efforts and ensuring timely interventions in challenging field environments.

85

## 86 Results

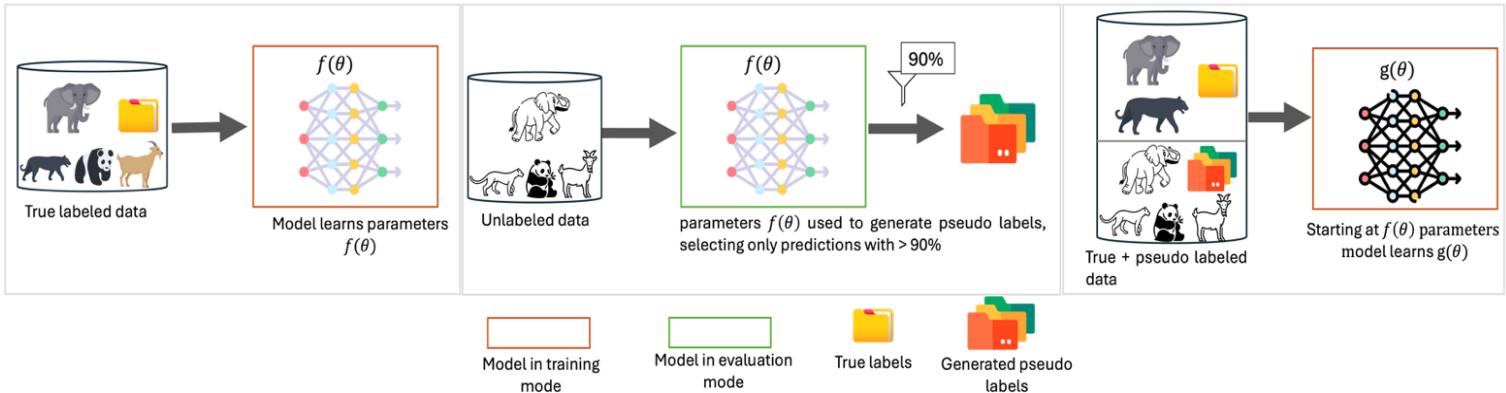
87 QuadPose employs a three-phase training strategy designed to improve pose estimation across  
88 quadrupeds. First, a data standardization step maps input data into two categories: elephants and  
89 other quadrupeds. Phase 1 uses a curriculum training approach on manually annotated data. Phase  
90 2 generates high-confidence pseudo-labels based on the model’s initial predictions. Finally, Phase  
91 3 applies a typical concurrent training approach, integrating both manually annotated and pseudo-  
92 labeled data for refinement (Fig. 1a). Below, we detail the datasets used to train QuadPose.

93

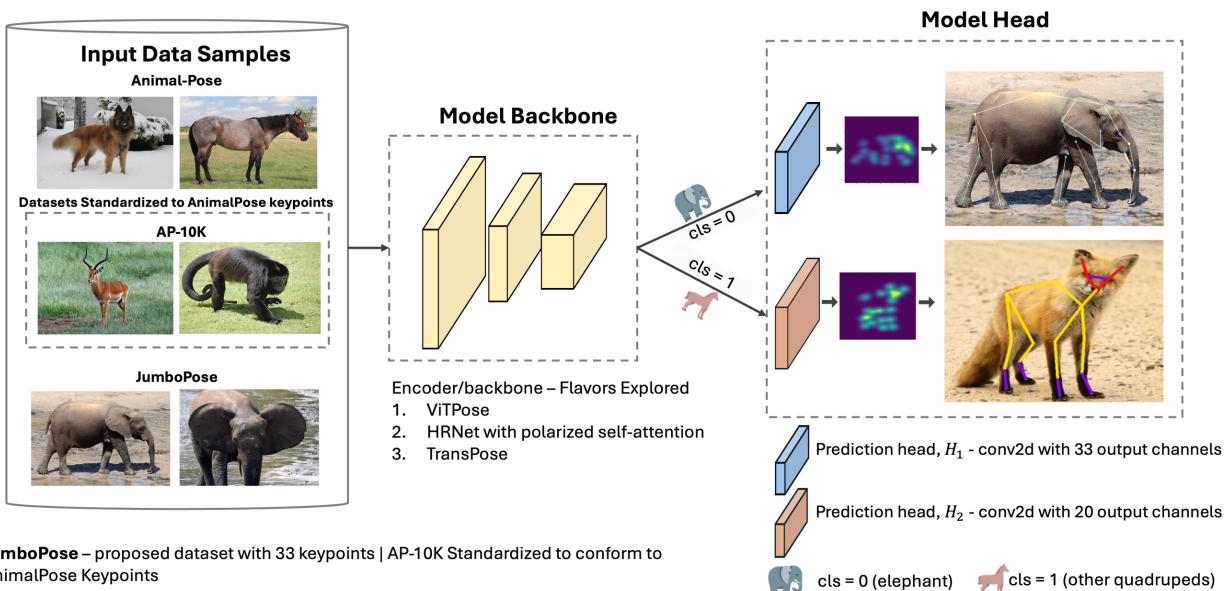
Phase 1: teacher model training

Phase 2: pseudo label generation

Phase 3: student model training

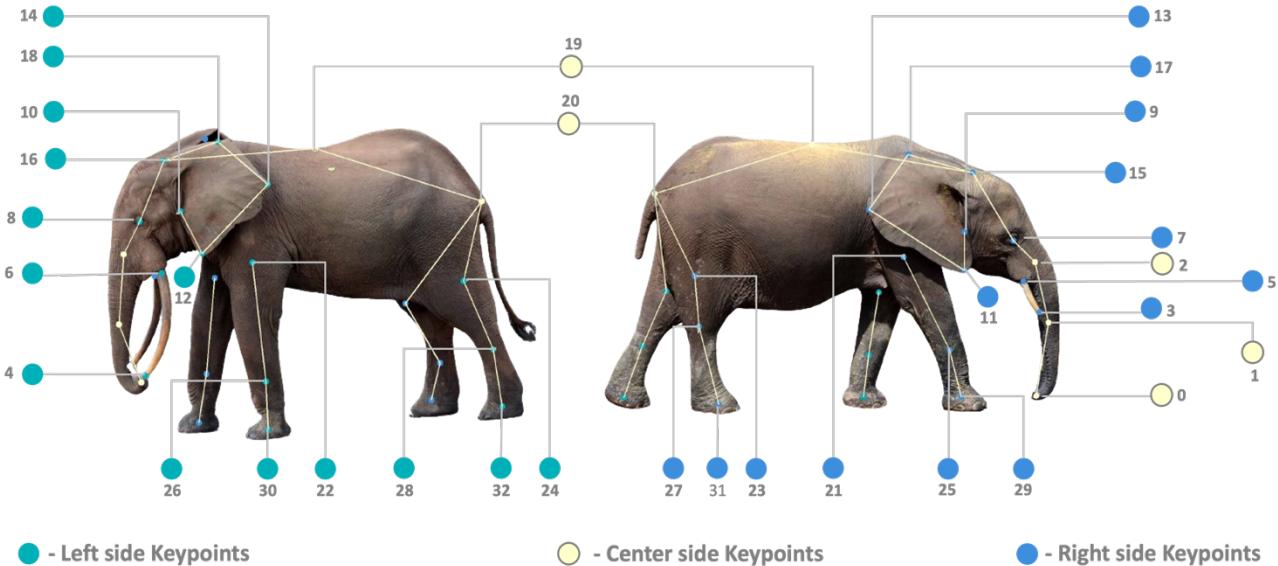


b



**JumboPose** – proposed dataset with 33 keypoints | AP-10K Standardized to conform to AnimalPose Keypoints

c



Keypoint	Definition	Keypoint	Definition	Keypoint	Definition	Keypoint	Definition
0	Bottom trunk	9	Right bottom ear	18	Top left tip ear	27	Right back knee
1	Mid trunk	10	Left bottom ear	19	Withers	28	Left back knee
2	Top trunk	11	Right bottom tip ear	20	Tail	29	Right front foot
3	Bottom right tusk	12	Left bottom tip ear	21	Right front elbow	30	Left front foot
4	Bottom left tusk	13	Right side tip ear	22	Left front elbow	31	Right back foot
5	Top right tusk	14	Left side tip ear	23	Right back elbow	32	Left back foot
6	Top left tusk	15	Top right ear	24	Left back elbow		
7	Right eye	16	Top left ear	25	Right front knee		
8	Left eye	17	Top right tip ear	26	Left front knee		

94  
95 Figure 1 Overview of the QuadPose framework. **a.** Schematic representation of the three-phase training  
96 pipeline. The teacher model is first trained using manually labeled data (Phase 1), then the learned  
97 parameters are used to generate high confidence pseudo labels from unlabeled data (Phase 2), and finally,  
98 a student model is trained using both manually and pseudo labeled data. **b.** Architecture of the QuadPose  
99 model, which is based on ViTPose<sup>40</sup>, HRNet<sup>39</sup> and, TransPose<sup>41</sup> model backbones. The model head is  
100 modified into dual prediction heads, first head is a convolutional layer that outputs 33 joints and second  
101 head is a convolutional layer that outputs 20 joints. **c.** Proposed Elephant Dataset, JumboPose, manually  
102 labeled dataset of 2,078 African elephants with 33 keypoints along with keys for the labelling scheme.  
103 Colored text highlight auxiliary keypoints unique to elephants.

## 104 Training Data

105 The QuadPose framework (Fig. 1b), standardizes data into two formats: one for elephants  
106 (JumboPose, Fig. 1c) and another for other quadrupeds. JumboPose, adapted from the ELPephants  
107 dataset<sup>42</sup> was originally developed for elephant Re-identification, while the other quadruped  
108 dataset integrates annotations from over 50 species across multiple sources. In phase 1, the teacher  
109 model is trained using 18,212 images from Animal-Pose<sup>37</sup>, AP-10K<sup>38</sup>, and JumboPose. Phase 2  
110 expands training samples by generating pseudo-labels from 53,600 unlabeled images across eleven  
111 datasets (see Methods). Finally, Phase 3 utilizes over 71,000 images from the previous phases for  
112 the student model training. Model weights from the teacher (phase 1) and student (phase 3) are  
113 publicly released along with the annotations for JumboPose and can be found at  
114 <https://github.com/QuadPose>.

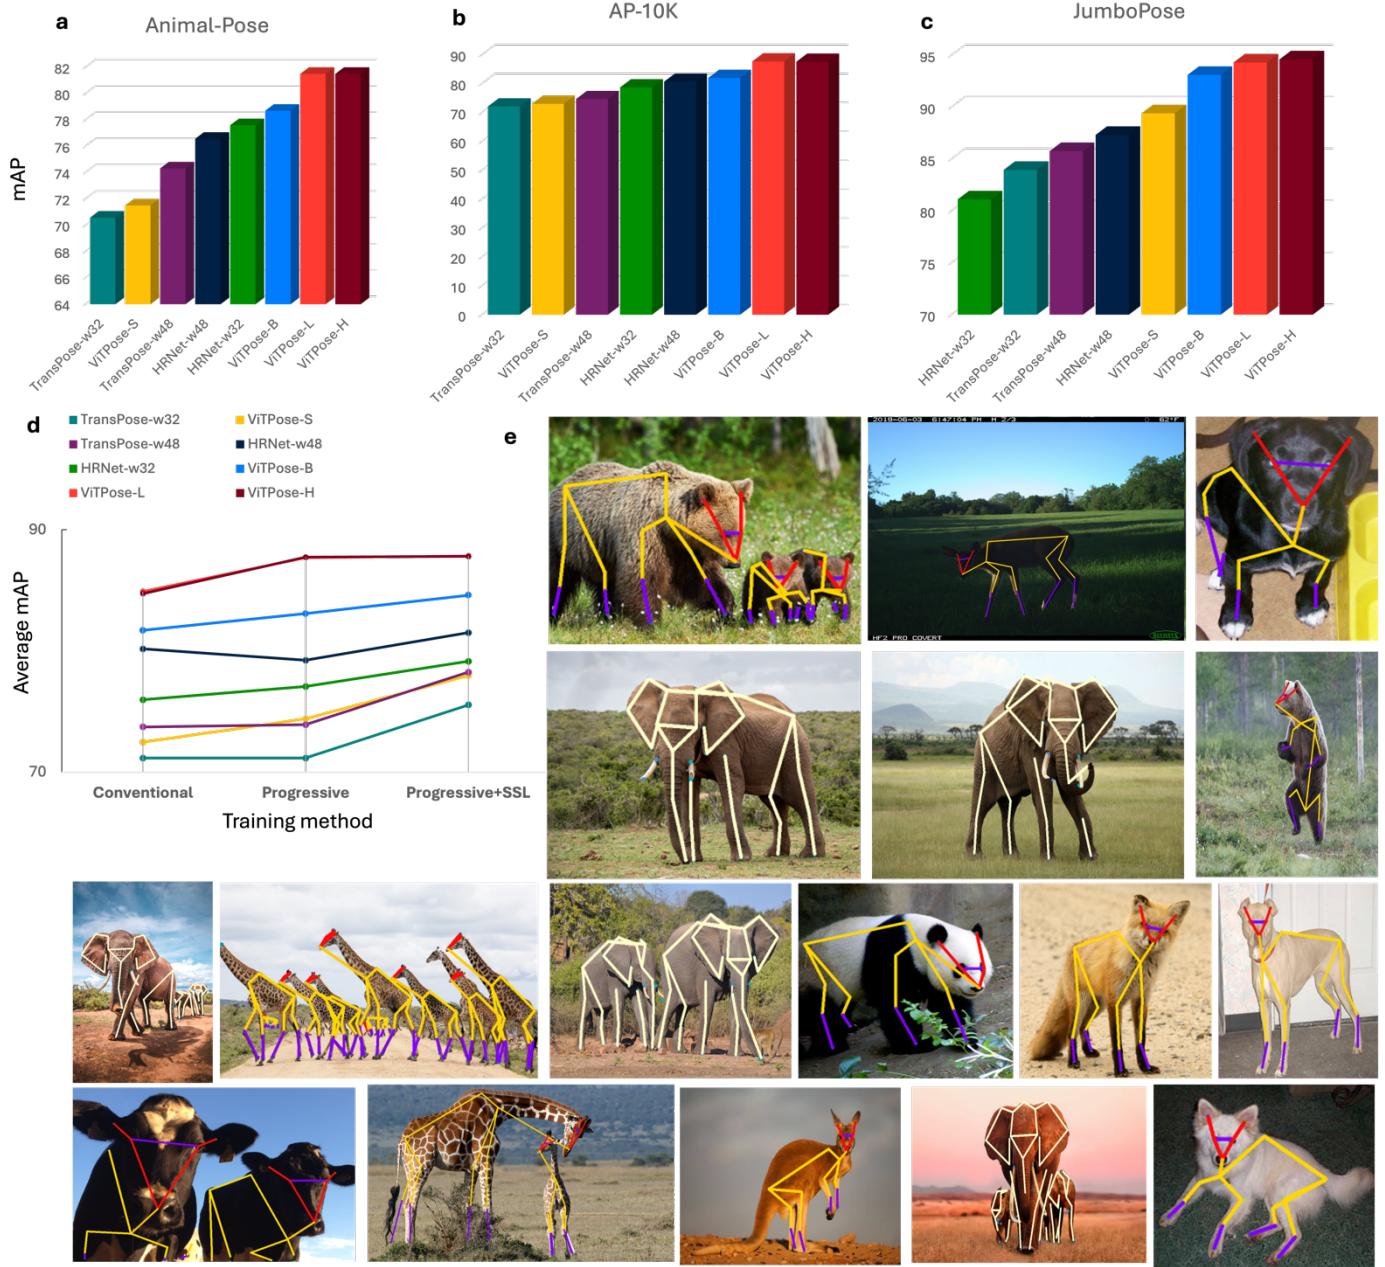
## 115 QuadPose method

116 The QuadPose is a robust method that handles diverse quadruped pose estimation datasets as two  
117 standardized datasets. It leverages top-down pose estimation<sup>40,43,39,41</sup> base architectures modified  
118 with dual prediction heads that conform to the defined unified standards. Top-down methods are  
119 chosen for their ability to first detect and classify objects before pose estimation, allowing for a  
120 binary classification step that dynamically routes input data to the appropriate prediction head (Fig.  
121 1b). In Phase 1, QuadPose follows a curriculum learning strategy, where training data is  
122 introduced progressively to improve model adaptability. By Phase 3, the model undergoes  
123 concurrent training, where both manually labeled and pseudo-labeled datasets are trained  
124 simultaneously to refine predictions across species.

## 125 Benchmarks

126 To evaluate the effectiveness of the QuadPose framework, we conducted experiments on three  
127 animal pose estimation datasets: Animal-Pose, AP-10K, and JumboPose. We compared three  
128 training strategies across multiple state-of-the-art architectures modified with dual prediction  
129 heads. The training strategies evaluated include 1) the Conventional training strategy, where all  
130 datasets are trained simultaneously, serving as a baseline, 2) Progressive training, where data is  
131 introduced incrementally using a curriculum learning approach, and 3) Progressive + SSL (semi-

132 supervised learning), which improves model generalization by using pseudo-labels. Fig. 2 and  
 133 Table 1 present the mean Average Precision (mAP) scores for each method and model  
 134 configuration. The global average mAP (Avg. mAP) and relative improvements ( $\Delta$  Avg. mAP)  
 135 relative to the conventional baseline method are also reported.



136 Figure 2. Benchmarking performance of QuadPose across datasets and training strategies. **(a-c)** Mean  
 137 Average Precision (mAP) scores for QuadPose on Animal-Pose, AP-10K, and JumboPose validation  
 138 datasets, evaluated across different architectures on the progressive + SSL training approach. The ViTPose  
 139 family of models achieves the highest mAP across all datasets. **d.** Comparative performance of the three  
 140 training strategies: Conventional baseline training, Progressive training, and Progressive + semi-supervised  
 141 learning (SSL) using pseudo-labels **e.** Qualitative results of the multi-task pose estimation problem, showing  
 142 keypoint predictions for elephants and other quadrupeds, the images used for qualitative analysis were  
 143 obtained from <sup>38,30,44</sup> and the Cummings School of Veterinary Medicine.

144 **Table 1 results:** Performance comparison of QuadPose across different architectures and training  
 145 strategies

Training Method	Architecture	Dataset (mAP)↑				
		Animal-Pose	AP-10K	JumboPose	Avg. mAP	Δ Avg. mAP
Conventional	HRNet-w48	73.90	74.8	<b>91.8</b>	80.17	-
Progressive	HRNet-w48	70.66	77.04	90.04	79.23	-0.94 ▼
Progressive + SSL	HRNet-w48	<b>76.57</b>	<b>80.71</b>	87.32	<b>81.5</b>	1.33 ▲
Conventional	HRNet-w32	68.49	69.43	<b>90.0</b>	75.97	-
Progressive	HRNet-w32	70.62	76.39	84.22	77.07	1.10 ▲
Progressive + SSL	HRNet-w32	<b>77.60</b>	<b>78.70</b>	81.11	<b>79.14</b>	3.17 ▲
Conventional	ViTPose-S	64.10	67.7	85.6	72.47	-
Progressive	ViTPose-S	66.50	68	88.7	74.40	1.93 ▲
Progressive + SSL	ViTPose-S	<b>71.50</b>	<b>73.00</b>	<b>89.40</b>	<b>77.97</b>	5.50 ▲
Conventional	ViTPose-B	76.40	78.5	90.2	81.70	-
Progressive	ViTPose-B	76.30	80.3	92.6	83.07	1.37 ▲
Progressive + SSL	ViTPose-B	<b>78.70</b>	<b>82.0</b>	<b>93.1</b>	<b>84.60</b>	2.90 ▲
Conventional	ViTPose-L	78.30	83.7	92.7	84.90	-
Progressive	ViTPose-L	80.40	<b>88.2</b>	<b>94.5</b>	87.70	2.80 ▲
Progressive + SSL	ViTPose-L	<b>81.50</b>	87.7	94.3	<b>87.83</b>	2.93 ▲
Conventional	ViTPose-H	77.90	83.3	93	84.73	-
Progressive	ViTPose-H	80.30	<b>88.2</b>	<b>94.7</b>	<b>87.73</b>	3.00 ▲
Progressive + SSL	ViTPose-H	<b>81.00</b>	87.6	94.6	<b>87.73</b>	3.00 ▲
Conventional	TransPose-w48	65.90	63.6	<b>91.7</b>	73.73	-
Progressive	TransPose-w48	68.12	65.28	86.53	73.92	0.19 ▲
Progressive + SSL	TransPose-w48	<b>74.32</b>	<b>74.66</b>	85.76	<b>78.25</b>	4.52 ▲
Conventional	TransPose-w32	62.65	60.41	<b>90.41</b>	71.16	-
Progressive	TransPose-w32	65.22	64.25	84.05	71.16	0.00
Progressive + SSL	TransPose-w32	<b>70.58</b>	<b>72.11</b>	83.95	<b>75.55</b>	4.39 ▲

146 The mAP scores (%) for Animal-Pose, AP-10K, and JumboPose datasets using different training strategies:  
 147 Conventional training (baseline), Progressive training, and Progressive + SSL (semi-supervised learning)  
 148 with pseudo-labels. The global average mAP across all datasets (Avg. mAP) and the relative improvement  
 149 ( $\Delta$  Avg. mAP) compared to the conventional baseline are reported. Results are shown for multiple model  
 150 architectures, including HRNet, ViTPose, and TransPose. The best-performing strategy for each  
 151 architecture is highlighted in bold.

## 152 Conventional Training

153 This training strategy provides the baseline for this study. Overall, it can be observed that the larger  
 154 model variants outperform their smaller counterparts across all datasets considered, with the  
 155 ViTPose-L and ViTPose-H achieving the highest average mAP of 84.9 and 84.73, respectively. In  
 156 the HRNet models, the HRNet-w48 achieves an average mAP of 80.17, while the HRNet-w32  
 157 attains an average mAP of 75.97. Similarly, within the TransPose family of architectures,  
 158 TransPose-w32 and TransPose-w48 report average mAP scores of 71.16 and 73.73, respectively.  
 159 This reflects the influence of model size on performance in this training paradigm. These baseline  
 160 results establish a reference for assessing the performance gains introduced by the progressive and  
 161 the progressive + SSL training strategies.

162

## 163 Progressive Training

164 The progressive training strategy aims to improve model generalization by gradually introducing  
165 training data in stages. Compared to conventional training, which processes all datasets  
166 simultaneously, progressive training significantly improves mAP scores across Animal-Pose and  
167 AP-10K, as shown in Table 1 and Fig. 2. In the ViTPose architectures, there is an average mAP  
168 improvement of 2.275, with all three datasets reporting consistent gains. The HRNet and  
169 TransPose architectures show similar trends, where the progressive training slightly outperforms  
170 or is at par with the conventional training. While an upward trend can be observed holistically, it  
171 can be seen from Table 1 that JumboPose performance tends to slightly decline in the HRNet and  
172 TransPose architectures as it transitions from conventional to progressive training strategies. This  
173 can be attributed to a few factors, including the continual learning approach adopted in training  
174 these architectures, where Animal-Pose and AP-10K are introduced first and second, for the task  
175 of learning representations to predict poses of general quadrupeds. Lastly, JumboPose is  
176 introduced for the task of learning representations to predict poses of elephants (see Methods).  
177 With this progressive training configuration, the model gradually catches up to learn elephant-  
178 specific poses; additionally, the sheer difference in volume between datasets for other quadrupeds  
179 and elephants is also a contributing factor (see methods, Table 2).

### 180 **Progressive + Semi Supervised Learning (SSL) Training**

181 Here a student model integrates pseudo-labels generated by a progressively trained teacher model,  
182 achieves the highest performance gains across all architectures. Notably, some models reach an  
183 increase of up to 5.5 in average mAP, demonstrating the effectiveness of this strategy.

184 For the ViTPose models, the ViTPose-S flavor sees the greatest performance boost from the  
185 progressive + SSL, with a 5.5 average mAP improvement over the baseline. Comparing the  
186 progressive (which serves as the teacher model) and the progressive + SSL (which serves as the  
187 student model) strategies, a 3.57 average mAP gain is recorded. For ViTPose-B variant, a 2.90  
188 average mAP increase is reported between the conventional baseline and progressive + SSL  
189 approaches, while a 1.56 average mAP improvement is observed between the progressive and  
190 progressive + SSL training approaches. For the larger variants, ViTPose-L and ViTPose-H, the  
191 gaps between the conventional and progressive + SSL training are 2.93 and 3.00 average mAP,  
192 respectively in favor of the progressive + SSL approach. However, analyzing the progressive  
193 (teacher model) and progressive + SSL, the performance gains while present begin to plateau, with  
194 marginal improvements of 0.18 on ViTPose-L and no improvement on the ViTPose-H. This  
195 suggests that while the Progressive + SSL is highly effective for smaller ViTPose architectures, its  
196 impact diminishes as model size increases, possibly due to the larger models already capturing rich  
197 feature representations during progressive training.

198 Similarly, the HRNet-based models enjoy a performance jump with the progressive + SSL  
199 approach. For HRNet-w48, which achieves an average mAP of 81.50, there is a 1.33 average mAP  
200 increase over the baseline. For HRNet-w32 with an average mAP of 79.14, there is a 3.17 average  
201 mAP increase over the baseline. The TransPose architectures also, see a significant performance  
202 boost with the progressive + SSL approach, whereas in the TrasnPose-w48, with an average mAP  
203 of 78.25, there is an average mAP increase of 4.52 when compared to the baseline. For the  
204 TransPose-w32 of average mAP 75.55, there is a 4.39 mAP gap over the baseline.

205

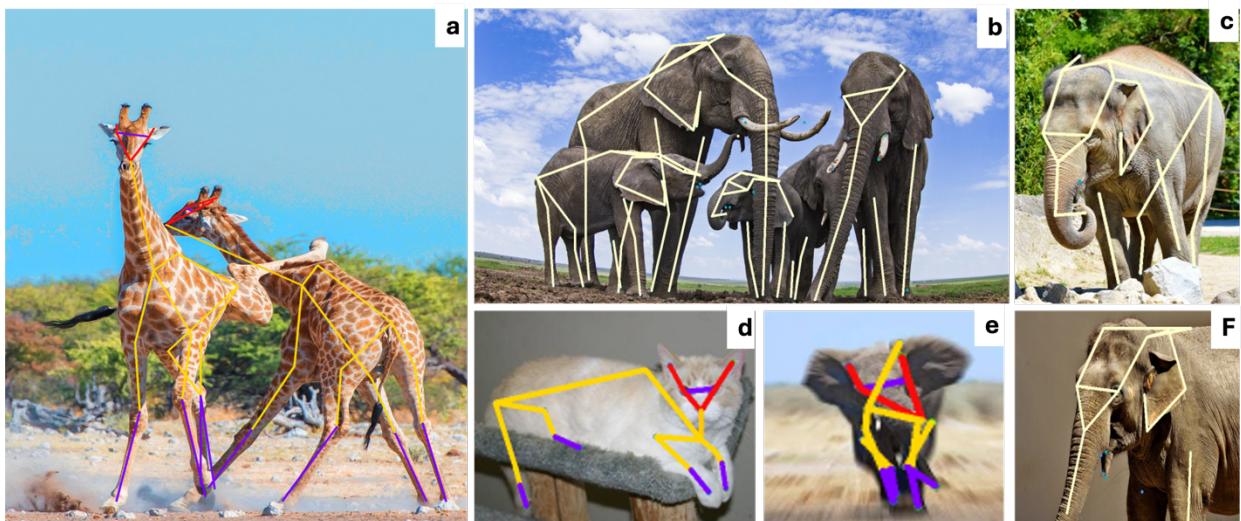
### 206 **Discussion**

207 This paper introduces QuadPose, a unified standard for pose estimation in mammals and a  
208 straightforward yet efficient framework for effectively utilizing noisy labels to improve a model’s  
209 generalizability to unseen animal breeds and species. By leveraging large amounts of unlabeled

210 data to generate pseudo-labels, the student model learns better feature representations for  
 211 localizing joints. Results show that this semi-supervised framework leads to an average  
 212 performance boost in mAP of about 3.5 across various state-of-the-art pose estimation models.  
 213 Furthermore, this work conducts a detailed performance evaluation of standard convolutional  
 214 neural networks and transformer-based networks for animal pose estimation. Results show that  
 215 ViTPose-based architectures perform best, as they demonstrate superior generalization capabilities  
 216 and robustness against catastrophic forgetting. Finally, this study presents JumboPose, a large-  
 217 scale dataset for elephant pose estimation and landmark localization.

218 Despite its strong performance improvements, QuadPose has certain limitations. One of its  
 219 primary challenges is handling complex and crowded scenes. Top-down models struggle in highly  
 220 occluded and densely populated scenarios typically due to detector performance where it might  
 221 miss objects that are heavily occluded or merge multiple objects of interest into a single bounding  
 222 box, in such detection failure cases there is no recourse to recovery<sup>37</sup>. Addressing this challenge  
 223 will require improved multi-instance pose tracking techniques and enhanced keypoint association  
 224 strategies. Additionally, since QuadPose relies on a pretrained object detector, classification errors  
 225 of the detector can further impact performance. Misclassification may lead to incorrect routing.  
 226 For example, if an elephant is wrongly detected as another animal, it will be routed to the general  
 227 quadruped prediction head instead of the elephant-specific head, and vice versa. This  
 228 misclassification introduces significant errors in keypoint localization and affects model  
 229 performance. Future improvements may involve integrating an uncertainty-aware detection system  
 230 or implementing self-correcting mechanisms to mitigate routing errors. Generalization to Asian  
 231 elephants is another limitation. The model is trained primarily on African elephants, which have  
 232 larger ears compared to their Asian counterparts. As a result, the model struggles to accurately  
 233 predict features of Asian elephants due to their distinct morphological differences. Expanding the  
 234 JumboPose to include Asian elephants will be crucial for improving the model's adaptability to  
 235 different elephant species. These limitations are shown in Fig. 3.

236



237

238

239 Figure 3. Examples of suboptimal results due to model limitations **a., b., d.** show the complex, crowded,  
 240 and occluded poses. **c. and f.** illustrate the output of the current model on Asian elephants with smaller  
 241 ears than their African counterparts **e.** error caused by the detector model misclassifying an elephant,  
 242 leading to it being incorrectly routed to the general quadruped prediction head.

243 Future research will focus on mitigating these challenges by reducing the framework's  
244 dependence on detector classification performance to minimize misrouting errors. Additionally,  
245 expanding the dataset to include Asian elephants will help improve species-specific generalization.  
246 Furthermore, developing models more robust to handling extremely crowded scenes will ensure  
247 better performance in real-world multi-animal environment where these animals often move in  
248 groups. By addressing these limitations, QuadPose can further enhance its applicability in wildlife  
249 conservation, behavioral analysis, and veterinary science, making it a more reliable solution for  
250 scalable and adaptable animal pose estimation.

251

## 252 **Methods**

### 253 **Datasets**

254 Three animal pose estimation datasets were used for phase 1 training and model performance  
255 evaluation: Animal-Pose<sup>37</sup>, AP-10K<sup>44</sup>, and JumboPose. Additionally, eleven supplementary  
256 datasets, which do not contain pose annotations, were used for pseudo-label generation. The  
257 following sections provide further details on these datasets.

258

#### 259 **Animal-Pose Dataset**

260 Cao et al.<sup>37</sup> proposed the Animal-pose dataset for developing cross-domain adaptation models for  
261 animal pose estimation. The dataset contains 6,117 annotated instances of cats, dogs, sheep, horses,  
262 and cows, with 20 keypoints defining their anatomical structures. For full details, refer to Cao et  
263 al.<sup>37</sup>.

264

#### 265 **AP-10K Dataset**

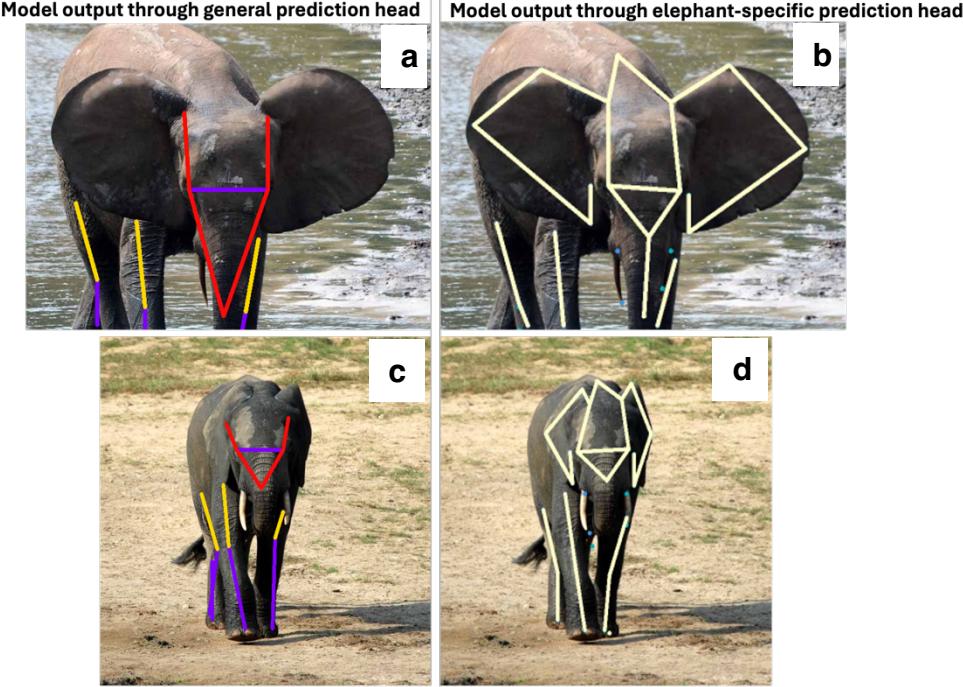
266 Yu et al.<sup>45</sup> curated the AP-10K, a large-scale manually annotated dataset for animal pose  
267 estimation. It consists of 10,015 images and over 13,000 labeled instances with 23 animal families  
268 and 54 species, making it the largest dataset for pose estimation in mammals. The keypoint  
269 annotation convention follows a structure similar to human pose estimation, with 17 keypoints  
270 representing anatomical features of the animals. For full details, refer to Yu et al.<sup>45</sup>.

271

#### 272 **JumboPose Dataset**

273 Elephants, particularly, African elephants possess distinctive features, such as large ears, trunks,  
274 and tusks, which differentiate them from most other quadrupeds. These features provide rich  
275 additional information that can enhance behavioral understanding. However, current annotation  
276 styles in state-of-the-art animal pose datasets do not capture these key anatomical distinctions.

277 To address this, we introduce JumboPose, a large-scale dataset specifically designed for elephant  
278 pose estimation. JumboPose contains over 2,000 labeled images and 6,617 unlabeled images of  
279 elephants. It consists of 33 keypoints with 20 keypoints defining the anatomical structure of the  
280 elephant similar to other quadrupeds in Animal-Pose<sup>37</sup>, and an additional 13 keypoints that capture  
281 the nuanced features specific to elephants. The elephant images annotated to create JumboPose  
282 was derived from the ElPephants dataset<sup>42</sup> originally designed for elephant re-identification. Fig.  
283 1c illustrates the annotation style used in JumboPose, while Fig. 4 highlights its motivation and  
284 structure.



285  
286  
287  
288  
289  
290

Figure 4. Comparison of model predictions using general quadruped vs. elephant-specific prediction head. (a., c.) Model outputs from the general quadruped prediction head, which fails to capture elephant-specific features such as trunks and ears. (b., d.) Model outputs from the elephant-specific prediction head, which correctly identifies distinctive anatomical features. These results highlight the importance of JumboPose in improving elephant pose estimation.

291  
292  
293  
294  
295  
296  
297  
298

### Unlabeled Data Curation

The unlabeled data used to generate pseudo labels were sourced from the following datasets; Macaque<sup>35</sup>, Tiger<sup>33</sup>, MS-COCO<sup>30</sup>, African Wildlife<sup>46</sup>, Animal Image Dataset<sup>47</sup>, Animals\_5<sup>30</sup>, Animals-10<sup>49</sup>, Endangered Animals<sup>50</sup>, IUCN Animals Dataset<sup>51</sup>, Wild Cats<sup>52</sup>, and Asian vs African Elephants<sup>53</sup>. Since MS-COCO is not explicitly an animal dataset, unlabeled elephant images were extracted by running it through an object detection algorithm (YOLOv8<sup>2</sup>).

**Table 2:** Dataset summarization.

Dataset	Species	Number of keypoints	Number of images
<b>Manually Labeled Data</b>			
Animal-Pose Dataset <sup>37</sup>	5	20	6,117
Ap-10K <sup>45</sup>	54	17	10,015
JumboPose (ours)	1	33	2,078
<b>Unlabeled Data for Pseudo-label Generation</b>			
Tiger <sup>33</sup>	1	-	4,124
Macaque <sup>35</sup>	1	-	13,085
MS COCO <sup>30</sup> (elephants)	-	-	2,202
African Wildlife <sup>46</sup>	4	-	1,504
Animal Image Dataset <sup>47</sup>	3	-	3,000
Animals_5 <sup>48</sup>	10	-	5,233
Animals-10 <sup>49</sup>	6	-	16,148
Endangered Animals <sup>50</sup>	4	-	800
IUCN Animals Dataset <sup>51</sup>	4	-	2,327
Wild Cats <sup>52</sup>	5	-	3,080

Asian vs African Elephants <sup>53</sup>	2	-	2,123
Total Pseudo labels - Elephants			6,617
Total Pseudo-labels - Other Quadrupeds			47,007
Total			71,834

299 Provides an overview of the datasets utilized for teacher model training (Phase 1) and pseudo-label  
300 generation (Phase 3). Three datasets (Animal-Pose, AP-10K, and JumboPose) were used for Phase 1  
301 training, comprising a total of 18,212 manually labeled images. For Phase 3, pseudo-labels were generated  
302 from the listed unlabeled datasets, contributing to a total of 71,834 images used in training.  
303 Note: The MS COCO dataset contains 2,202 elephant images, extracted specifically for this study.

### 304 Dataset Standardization

305 The data is standardized into two formats based on quadruped type: elephants and other  
306 quadrupeds. The elephant type is straightforward, as JumboPose is the only dataset explicitly  
307 designed for elephant pose estimation, to the best of our knowledge. For other quadrupeds,  
308 differences in the number of keypoints, labeling style, and keypoint location across datasets  
309 necessitate a standardized format for model training. The anatomical structure proposed by Cao et  
310 al.<sup>37</sup>, is adopted as it accounts for an additional number of useful keypoints like the ears and  
311 withers. Keypoints that are not annotated in certain datasets are assigned a "missing" tag and left  
312 unannotated to maintain consistency across datasets. Fig. 5 shows how the AP-10K annotation  
313 style is remapped accordingly.

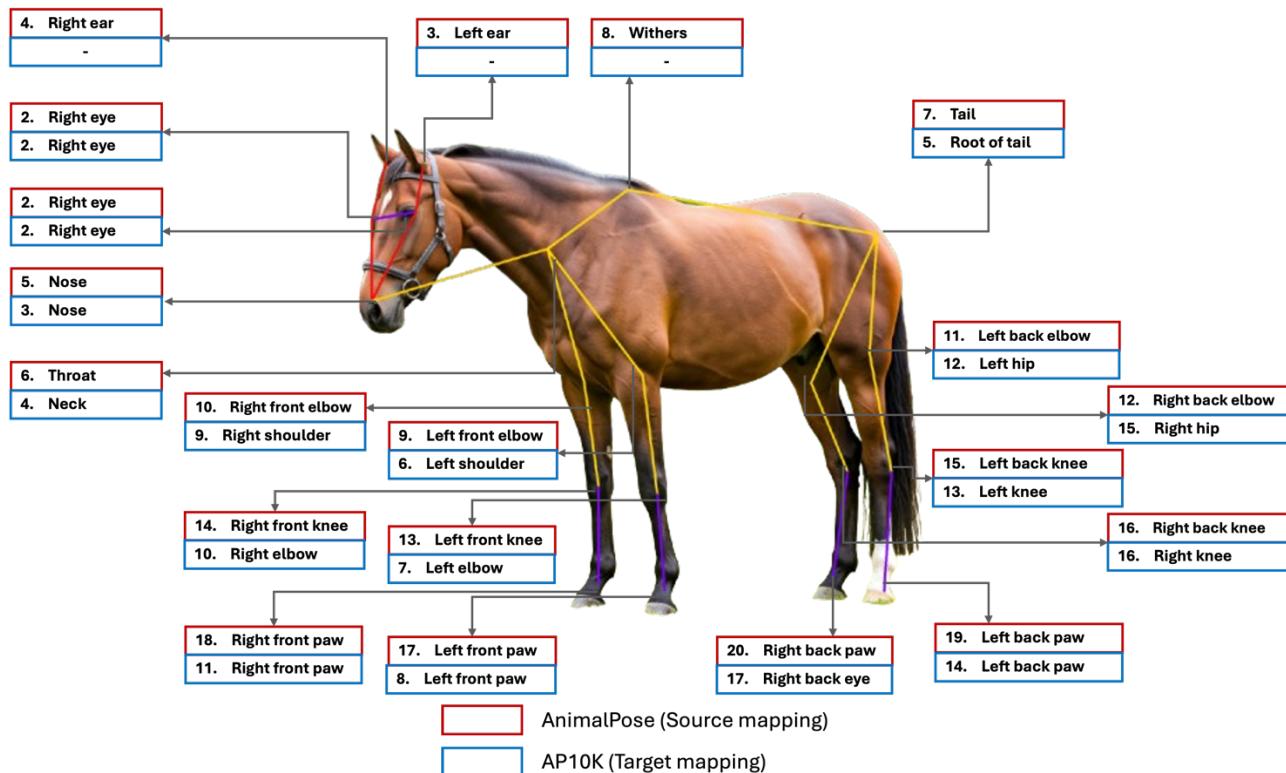


Figure 5. This figure illustrates how AP-10K keypoints are remapped and standardized to align with the Animal-Pose anatomical structure.

314

### 315 QuadPose Training

The proposed model is tasked with two objectives: 1) elephant pose estimation, routed to prediction head,  $H_1$  and 2) other quadruped pose estimation, routed to prediction head,  $H_2$ . The general architectural overview is provided in Fig. 1b. The dual-head architecture gradually learns and retains good representations for both tasks. As shown in Fig. 1a, there are three phases in the training pipeline. Curriculum learning, which is a method used to gradually increase the complexity of data samples in the training process<sup>54</sup> is adopted in training, specifically for task 2. A simple yet efficient curriculum learning strategy is adopted, where data samples are introduced in order of difficulty, progressing from easy examples (examples with well-defined poses with most joints visible) to more challenging ones (examples of animals with occluded or missing joints). This strategy is adopted in training the teacher ( $f_\theta$ ) network.

For the first phase, training is done on three datasets: JumboPose (33 keypoints) for task 1, and Animal pose (20 keypoints) and AP-10k (17 keypoints) for task 2. A curriculum based on the number of annotated keypoints is adopted. This helps determine which datasets are prioritized per task at multiple intervals during training, ensuring the model progressively learns from data distributions that provide the most information before adapting to less informative ones, traversing from the known to the less known. This structured progression improves cross-dataset learning and enhances the model’s generalization capabilities across different species. The training strategies empirically derived to give the best performance by model architecture are:

a) ViTPose-based models:

- Initial training begins with JumboPose and Animal-Pose as it simultaneously provides standard baselines for both tasks.
- Finally, AP-10K dataset is progressively introduced allowing for better refinement of task 2 representations

b) HRNet and TransPose-based models

- Initial training starts with Animal-Pose as it provides the standard baseline for task 2.
- AP-10k is introduced.
- Finally, JumboPose is introduced in a continual learning manner to learn task 1 representations while retaining task 2’s representations.

In the second phase, pseudo-labels are generated from a large set of unlabeled data of unseen animal species using the learned weights of the teacher network  $f_\theta$ . Each image sample is classified as either ‘easy’ or ‘hard’ based on the pose complexity, which is determined by the number of high-confidence keypoints detected in each instance. The underlying hypothesis is that an animal’s pose is better defined when its complete anatomical structure is visible. By setting a high confidence threshold, unreliable labels are filtered out, ensuring only accurate labels contribute to training. Due to resource constraints and to prevent the models from over-committing to low-confidence and potentially misleading labels, only ‘easy’ samples are utilized in this study.

In the third phase, the ground-truth annotated datasets, and the pseudo-labeled data are used to train the student network,  $g_\theta$ . The teacher network distills learned knowledge into the student, improving the model’s ability to generalize across diverse species. This knowledge distillation process helps the student network capture robust feature representations and enhances overall pose estimation.

### Model Specific Training Configurations

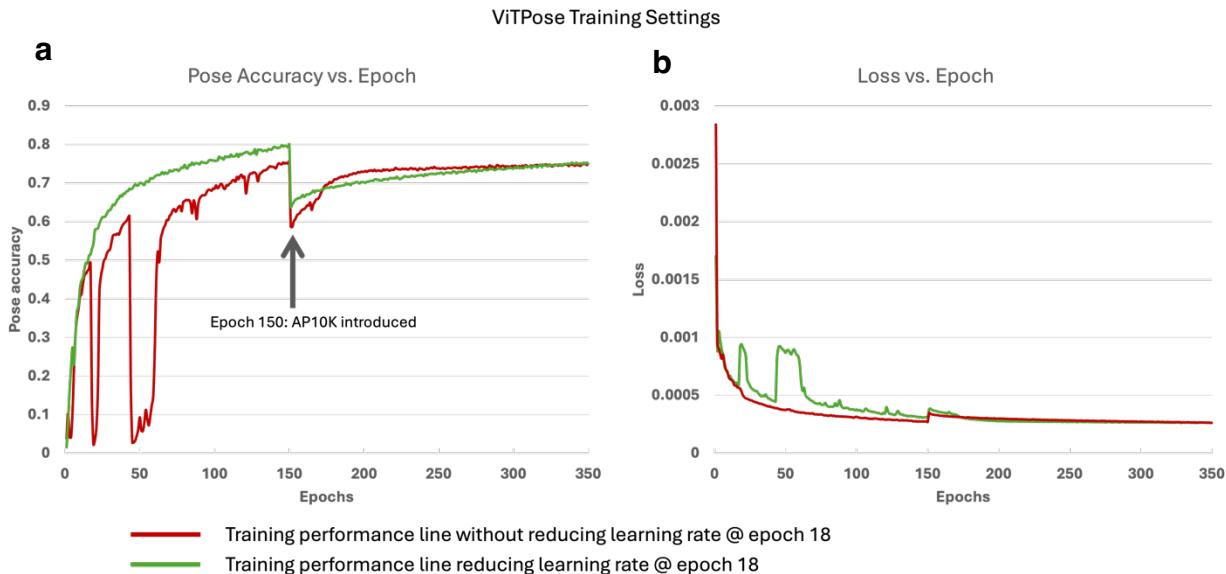
The following shows the specific learning settings employed for each architecture. All models were trained on a single NVIDIA V100 GPU and implemented using the PyTorch deep learning framework<sup>55</sup>.

361 **ViTPose model training.** The learning rate was reduced at epochs 18 and 150 to improve training  
 362 stability and performance. Initially, the learning rate reduction was scheduled only at epoch 150,  
 363 coinciding with the introduction of the AP-10K dataset. However, empirical analysis revealed that  
 364 an earlier reduction at epoch 18 further stabilizes training and enhances convergence. As shown  
 365 in Fig. 6, early reduction in learning rate improves training stability. Table 3 summarizes the  
 366 training settings used for the ViTPose family of models.  
 367

**Table 3:** ViTPose Training Settings

Starting LR	0.0005
LR Scheduler	Multi-step
LR Step Factor	0.1
LR Steps (epochs)	[18, 150, 400, 450]
Batch size	32
Optimizer	Adam <sup>56</sup>
<b>Data Arrangement for progressive training Phase 1</b>	
Epochs 0 - 149	Animal-Pose + JumboPose
Epochs 150 - 349	Animal-Pose + AP-10K + JumboPose
<b>Data Arrangement for progressive training Phase 3</b>	
Epochs 350 - 500	Animal-Pose + AP-10K + Pseudo_other_quadrupeds + JumboPose + Pseudo_elephants

368



369 **Figure 6.** ViTPose-B phase 1 training performance **a.** The pose accuracy per epoch **b.** The Loss per epoch:  
 370 the red line - when the learning rate is not reduced at the 18<sup>th</sup> epoch, leading to an unstable training that  
 371 eventually converges, the green line - when the learning rate is reduced at the 18<sup>th</sup> epoch leading to a more  
 372 stable training with better performance. Note, the dip at the 150<sup>th</sup> epoch occurs as a result of introducing  
 373 AP-10K dataset in the progressive training.  
 374

375 **HRNet and TransPose model training.** These models were found to be more sensitive to training  
 376 with the proposed dual-prediction head modification compared to the ViTPose-based models.  
 377 Hence, the need for a carefully designed continual learning approach. A common issue  
 378 encountered in continual learning is catastrophic forgetting, where models lose previously learned  
 379 information when adapting to new tasks<sup>57</sup>. This occurs because the model continuously updates  
 380

its parameters, often overwriting past knowledge from an initial task in favor of a more recent task. An ideal multitask model mitigates this by learning an optimal parameter space that effectively generalizes across tasks while preserving previously acquired knowledge. To mitigate catastrophic forgetting in HRNet and TransPose-based models, we adopted the training settings outlined in Table 4. Specifically, we reduced the learning rate by a smaller multiplication factor of 0.05, which helps in stabilizing updates and retaining existing knowledge. By implementing a more gradual learning rate reduction, the model's updates become less aggressive, thereby minimizing the risk of overwriting previously learned information.

**Table 4:** HRNet & TransPose Training Settings

Starting LR	0.001
LR Scheduler	Multi-step
LR Step Factor	0.05
LR Steps (epochs)	120
Batch size (HRNet   TransPose)	32   24
Optimizer	Adam <sup>56</sup>
<b>Data Arrangement for progressive training Phase 1</b>	
Epochs 0 - 74	Animal-Pose
Epochs 75 - 119	Animal-Pose + AP-10K
Epochs 120 - 349	Animal-Pose + AP-10K + JumboPose
<b>Data Arrangement for progressive training Phase 3</b>	
Epochs 350 - 500	Animal-Pose + AP-10K + Pseudo_other_quadrupeds + JumboPose + Pseudo_elephants

### Cost function

The mean squared error (MSE) is used to evaluate the Euclidian distance between the predicted and ground-truth keypoints. The objective function,  $L_{joint}$  for training the network using a set of labeled and pseudo-labeled datasets, is given as

$$L_{joint} = L_{sup} + cL_{pseudo} \quad (1)$$

Where  $L_{sup}$  and  $L_{pseudo}$  are the mean squared errors (MSE) on the manually labeled data (supervised) and pseudo-labeled examples, respectively. The binary flag  $c$  determines if pseudo-labels are used:

$$c = \begin{cases} 1, & \text{if using pseudo-labels (student training, phase 3)} \\ 0, & \text{otherwise (teacher training, phase 1)} \end{cases} \quad (2)$$

Each loss term,  $L_{sup}$  or  $L_{pseudo}$  is the sum of MSE losses for task 1 (elephant pose estimation) and task 2 (other quadruped pose estimation):

$$L = L_{MSE}(y_i^1, \hat{y}_i^1) + L_{MSE}(y_i^2, \hat{y}_i^2) \quad (3)$$

Where  $y_i^1$  and  $y_i^2$  are the ground-truth labels for tasks 1 and 2, respectively while  $\hat{y}_i^1$  and  $\hat{y}_i^2$  correspond to model predictions for tasks 1 and 2.

$$L_{pseudo} = \sum_{j=0}^N w_j^\emptyset L(I_j^T, m(I_j^T | \emptyset)) \quad (4)$$

409  
410  $L_{pseudo}$  is similar to equation (8) in the work proposed by Cao et al.<sup>37</sup> for self-paced selection of  
411 pseudo-labels and  $m(I_j^T | \emptyset)$  is the output by the model of current weights  $\emptyset$  on an input image  $I_j^T$ .  
412  $w_j^\emptyset$  denotes whether the pose prediction on  $I_j^T$  is a hard or easy example.

$$w_j^\emptyset = \begin{cases} 1.0, & \text{if } C(m(I_j^T | \emptyset)) > \mu \\ 0, & \text{otherwise} \end{cases} \quad (5)$$

413  
414 Where  $C(m(I_j^T | \emptyset))$  denotes the output confidence score on  $I_j^T$  by the current model and  $\mu$  is the  
415 threshold to filter unreliable outputs.  $\mu$  is set to 0.9 for our task to select only high confidence  
416 pseudo-labels.

#### 417 Algorithms

418 **Algorithm 1** outlines the training process using the dual prediction head architecture. Where a  
419 data sample could either be elephant ( $k = 1$ ) or other quadrupeds ( $k = 2$ ). Both sample types pass  
420 through the same encoder backbone, which extracts features and learn representations common to  
421 all quadrupeds. For the prediction heads, if  $k = 1$ (elephant), the extracted features are routed to the  
422 prediction head,  $H_1$  to generate 33 heatmaps corresponding to the keypoints. If  $k = 2$  (other  
423 quadrupeds), the extracted features are routed to prediction head,  $H_2$  to produce 20 heatmaps  
424 corresponding to the predicted key points.

---

#### Algorithm 1 Dual-Head Model Training

---

**Input:** Training data,  $D = (x_i^k, y_i^k) : k \in \{1, 2\}$   
 $x_i^k, y_i^k \triangleq i$ th input image and ground truth respectively  
Encoder backbone,  $Enc$   
Prediction Heads,  $H_1$  and  $H_2$

1: **Initialize** Encoder backbone  $Enc$  and Prediction Heads  $H_2$  and  $H_2$

2: **repeat**

- 3:  $(x_i^k, y_i^k) \sim D \triangleq$  sample batch from dataset
- $\emptyset = Enc(x_i^k) \triangleq$  data goes into encoder backbone
- 4: *if*  $k == 1$ :
- 5:      $\hat{y}_i^1 = H_1(\emptyset) \triangleq$  Route  $\emptyset$  to  $H_1$
- 6: *elif*  $k == 2$ :
- 7:      $\hat{y}_i^2 = H_2(\emptyset) \triangleq$  Route  $\emptyset$  to  $H_2$
- 8:      $L = loss(y_i^1, \hat{y}_i^1) + loss(y_i^2, \hat{y}_i^2)$
- 9:     Take gradient step to update  $Enc, H_1$  and  $H_2$

10: **until** training converges

---

425 **Algorithm 2** describes the inference process of the dual prediction head architecture. This strongly  
426 leverages the top-down nature of the pose estimation algorithms considered. First, an image is  
427 processed by a detector, which outputs the object class along with the bounding box coordinates.

428 The detected class is then binarized, assigning ‘1’ to elephants and ‘2’ to other quadrupeds. The  
 429 bounding box coordinates define a cropped image region, which is passed through the encoder  
 430 backbone for feature extraction. The extracted features are then routed identically to Algorithm 1.

---

## Algorithm 2 Dual-Head Model Inferencing

---

**Input:** Image =  $x$   
 Detector  $d$   
 Trained model Encoder backbone  $Enc$  and Prediction Heads  $H_1$  and  $H_2$

- 1: **Load trained** Encoder backbone,  $Enc$  and Prediction Heads,  $H_2$ , and  $H_2$  weights
- 2:  $x^k = d(x) \triangleleft$  detected object from image with predicted class,  $k$
- 3:  $\emptyset = Enc(x^k) \triangleleft$  data goes into encoder backbone
- 4: if  $k == 1$ : predicted class is ‘1’ (elephant class)
- 5:      $\hat{y}_i^1 = H_1(\emptyset) \triangleleft$  Route  $\emptyset$  to  $H_1$
- 6: elif  $k == 2$ : predicted class is ‘2’ (other quadruped class)
- 7:      $\hat{y}_j^2 = H_2(\emptyset) \triangleleft$  Route  $\emptyset$  to  $H_2$

---

431  
 432 **Model Architectures**  
 433 QuadPose employs top-down models, including ViTPose<sup>40</sup>, HRNet with polarized self-attention  
 434 for improved representation capacity<sup>39,18</sup> and TransPose<sup>41</sup>, each modified with dual prediction  
 435 heads to enable multi-task pose estimation for elephants and other quadrupeds. This method  
 436 follows a detection-first approach, which isolates objects of interest using bounding boxes while  
 437 simultaneously leveraging class information as a signal to route the detected objects to the  
 438 appropriate prediction head. Specifically, the classification is binarized into two categories: (1)  
 439 Class 1- if an elephant is detected, it is routed to prediction head  $H_1$  and (2) Class 2 – if a general  
 440 quadruped is detected, it is routed to prediction head  $H_2$ . This adaptive routing mechanism ensures  
 441 that each animal is processed by the most relevant prediction head. A pretrained faster R-CNN<sup>1</sup>  
 442 model is used for object detection during inference.

443 **Evaluation Metric**

444 The Object Keypoint Similarity (OKS)<sup>58</sup> calculates the distance between predicted keypoints, and  
 445 ground-truth points normalized by the object’s scale<sup>59</sup>. OKS values serve as thresholds for  
 446 computing mean average precision (mAP), which ranges from 0 to 1, with higher values indicating  
 447 more accurate keypoint localization.

448

$$OKS = \frac{\sum_i \exp(-d_i^2 / 2s^2 k_i^2) \delta(v_i > 0)}{\sum_i \delta(v_i > 0)} \quad (6)$$

449

450 Where,

- 451 •  $d_i$  is the Euclidean distance between the ground-truth and predicted keypoint  
 452 •  $s$  is the square root of the object segment area  
 453 •  $k$  is the per-keypoint constant that controls fall off  
 454 •  $v_i$  is the visibility flag that can be 0, 1, 2 for not labeled, labeled but not visible and visible  
 455 and labeled respectively  
 456 •  $\delta(v_i > 0)$  ensures only labeled keypoints contribute

457 **Acknowledgements**

458 The authors thank Allen Rutberg for his valuable insights and Victor Oludare for his insightful  
459 contributions and code assistance. Funding was partly provided by the Humane Society of the  
460 United States and the Cummings School of Veterinary Medicine at Tufts University, the Tufts  
461 Elephant Conservation Alliance (TECA), and The Bailey Wildlife Foundation.

462 **Author Contributions**

463 Conceptualization, K. Panetta, O. J., and S. R.; Methodology, O. J., S. R., and K. Panetta; Software,  
464 K. Panetta, O. J., S. R., J. H.; Validation, K. P., O. J., S. R., J. H., S. A.; Formal Analysis, K.  
465 Panetta, O. J., S. A.; Investigation, O. J., S. R., J. H.; K. Panetta.; Resources, K. Panetta.; Data  
466 Curation, O. J.; Writing—Original Draft, O. J., S. R., K. Panetta.; Writing—Review and Editing,  
467 K. Panetta, O. J., S. R., J. H., S. A., K. Pereira; Visualization, O. J., K. Pereira; Supervision, K.  
468 Panetta, S.A.; Project Administration, K. Panetta, S.A.; Funding Acquisition, K. Panetta. All  
469 authors have read and agreed to the published version of the manuscript.

470 **Competing Interests**

471 The authors declare no competing interests.

472 **Data Availability**

473 JumboPose is made publicly available with download instructions at <https://github.com/QuadPose>.

474 **Code Availability**

475 QuadPose source code is available at <https://github.com/QuadPose>. All other requests should be  
476 made to the corresponding author.

477 **References**

- 478 1. Ren, S., He, K., Girshick, R. & Sun, J. Faster R-CNN: Towards Real-Time Object Detection with Region Proposal Networks. *IEEE Trans. Pattern Anal. Mach. Intell.* **39**, 1137–1149 (2015).
- 479 2. Varghese, R. & M., S. YOLOv8: A Novel Object Detection Algorithm with Enhanced Performance and Robustness. in *2024 International Conference on Advances in Data Engineering and Intelligent Computing Systems (ADICS)* 1–6 (IEEE, Chennai, India, 2024). doi:10.1109/ADICS58448.2024.10533619.
- 480 3. Carion, N. *et al.* End-to-End Object Detection with Transformers. Preprint at <https://doi.org/10.48550/ARXIV.2005.12872> (2020).
- 481 4. Liu, W. *et al.* SSD: Single Shot MultiBox Detector. (2015) doi:10.48550/ARXIV.1512.02325.
- 482 5. He, K., Zhang, X., Ren, S. & Sun, J. Deep Residual Learning for Image Recognition. *Proc. IEEE Comput. Soc. Conf. Comput. Vis. Pattern Recognit.* **2016-December**, 770–778 (2015).
- 483 6. Ronneberger, O., Fischer, P. & Brox, T. U-Net: Convolutional Networks for Biomedical Image Segmentation. *Lect. Notes Comput. Sci. Subser. Lect. Notes Artif. Intell. Lect. Notes Bioinforma.* **9351**, 234–241 (2015).
- 484 7. He, K., Gkioxari, G., Dollár, P. & Girshick, R. Mask R-CNN. Preprint at <https://doi.org/10.48550/ARXIV.1703.06870> (2017).
- 485 8. Oludare, V., Kezebou, L., Jinadu, O., Panetta, K. & Agaian, S. Attention-based two-stream high-resolution networks for building damage assessment from satellite imagery. in *Multimodal Image Exploitation and Learning 2022* (eds. Agaian, S. S., Asari, V. K., DelMarco, S. P. & Jassim, S. A.) vol. 12100 121000L (SPIE, 2022).
- 486 9. Kirillov, A. *et al.* Segment Anything. Preprint at <https://doi.org/10.48550/ARXIV.2304.02643> (2023).
- 487 10. Chen, L.-C., Papandreou, G., Kokkinos, I., Murphy, K. & Yuille, A. L. DeepLab: Semantic Image Segmentation with Deep Convolutional Nets, Atrous Convolution, and Fully Connected CRFs. Preprint at <https://doi.org/10.48550/ARXIV.1606.00915> (2016).

- 501 11. Cao, Z., Hidalgo, G., Simon, T., Wei, S. E. & Sheikh, Y. OpenPose: Realtime Multi-Person 2D Pose Estimation  
502 using Part Affinity Fields. *IEEE Trans. Pattern Anal. Mach. Intell.* **43**, 172–186 (2018).
- 503 12. Mathis, A. *et al.* DeepLabCut: markerless pose estimation of user-defined body parts with deep learning. *Nat.*  
504 *Neurosci.* **21**, 1281–1289 (2018).
- 505 13. Ye, S. *et al.* SuperAnimal pretrained pose estimation models for behavioral analysis. *Nat. Commun.* **15**, 5165  
506 (2024).
- 507 14. Pereira, T. D. *et al.* Fast animal pose estimation using deep neural networks. *Nat. Methods* **16**, 117–125 (2019).
- 508 15. Lauer, J. *et al.* Multi-animal pose estimation, identification and tracking with DeepLabCut. *Nat. Methods* **19**, 496–  
509 504 (2022).
- 510 16. Dong, C. & Du, G. An enhanced real-time human pose estimation method based on modified YOLOv8 framework.  
511 *Sci. Rep.* **14**, 8012 (2024).
- 512 17. Yang, J., Zeng, A., Zhang, R. & Zhang, L. X-Pose: Detecting Any Keypoints. Preprint at  
513 <https://doi.org/10.48550/ARXIV.2310.08530> (2023).
- 514 18. Wang, J. *et al.* Deep High-Resolution Representation Learning for Visual Recognition. *IEEE Trans. Pattern Anal.*  
515 *Mach. Intell.* **43**, 3349–3364 (2021).
- 516 19. Vaswani, A. *et al.* Attention Is All You Need. Preprint at <https://doi.org/10.48550/ARXIV.1706.03762> (2017).
- 517 20. Dosovitskiy, A. *et al.* An Image is Worth 16x16 Words: Transformers for Image Recognition at Scale. Preprint at  
518 <https://doi.org/10.48550/ARXIV.2010.11929> (2020).
- 519 21. Ning, G., Liu, P., Fan, X. & Zhang, C. A Top-down Approach to Articulated Human Pose Estimation and Tracking.  
520 *Lect. Notes Comput. Sci. Subser. Lect. Notes Artif. Intell. Lect. Notes Bioinforma.* **11130 LNCS**, 227–234 (2019).
- 521 22. Geng, Z., Sun, K., Zhang, Z. & Wang, J. Bottom-Up Human Pose Estimation Via Disentangled Keypoint  
522 Regression.
- 523 23. Cao, Z., Simon, T., Wei, S. E. & Sheikh, Y. Realtime Multi-Person 2D Pose Estimation using Part Affinity Fields.  
524 *Proc. - 30th IEEE Conf. Comput. Vis. Pattern Recognit. CVPR 2017 2017-January*, 1302–1310 (2016).
- 525 24. Iqbal, U. & Gall, J. Multi-Person Pose Estimation with Local Joint-to-Person Associations. Preprint at  
526 <https://doi.org/10.48550/ARXIV.1608.08526> (2016).
- 527 25. Gkioxari, G., Hariharan, B., Girshick, R. & Malik, J. Using k-Poselets for Detecting People and Localizing Their  
528 Keypoints. in *2014 IEEE Conference on Computer Vision and Pattern Recognition* 3582–3589 (IEEE, Columbus,  
529 OH, USA, 2014). doi:10.1109/CVPR.2014.458.
- 530 26. Papandreou, G. *et al.* Towards Accurate Multi-person Pose Estimation in the Wild. Preprint at  
531 <https://doi.org/10.48550/ARXIV.1701.01779> (2017).
- 532 27. Li, M., Zhou, Z., Li, J. & Liu, X. Bottom-up Pose Estimation of Multiple Person with Bounding Box Constraint.
- 533 28. Pishchulin, L. *et al.* DeepCut: Joint Subset Partition and Labeling for Multi Person Pose Estimation. Preprint at  
534 <https://doi.org/10.48550/ARXIV.1511.06645> (2015).
- 535 29. Insafutdinov, E., Pishchulin, L., Andres, B., Andriluka, M. & Schiele, B. DeeperCut: A Deeper, Stronger, and  
536 Faster Multi-Person Pose Estimation Model. Preprint at <https://doi.org/10.48550/ARXIV.1605.03170> (2016).
- 537 30. Lin, T.-Y. *et al.* Microsoft COCO: Common Objects in Context. in *Computer Vision – ECCV 2014* (eds. Fleet, D.,  
538 Pajdla, T., Schiele, B. & Tuytelaars, T.) vol. 8693 740–755 (Springer International Publishing, Cham, 2014).
- 539 31. Zhang, S.-H. *et al.* Pose2Seg: Detection Free Human Instance Segmentation. in *2019 IEEE/CVF Conference on*  
540 *Computer Vision and Pattern Recognition (CVPR)* 889–898 (IEEE, Long Beach, CA, USA, 2019).  
541 doi:10.1109/CVPR.2019.00098.
- 542 32. Andriluka, M., Pishchulin, L., Gehler, P. & Schiele, B. 2D Human Pose Estimation: New Benchmark and State of  
543 the Art Analysis. in *2014 IEEE Conference on Computer Vision and Pattern Recognition* 3686–3693 (IEEE,  
544 Columbus, OH, USA, 2014). doi:10.1109/CVPR.2014.471.
- 545 33. Li, S., Li, J., Tang, H., Qian, R. & Lin, W. ATRW: A Benchmark for Amur Tiger Re-identification in the Wild.  
546 *MM 2020 - Proc. 28th ACM Int. Conf. Multimed.* **20**, 2590–2598 (2019).
- 547 34. Mathis, A. *et al.* Pretraining boosts out-of-domain robustness for pose estimation. in *2021 IEEE Winter Conference*  
548 *on Applications of Computer Vision (WACV)* 1858–1867 (IEEE, Waikoloa, HI, USA, 2021).  
549 doi:10.1109/WACV48630.2021.00190.

- 550 35. Labuguen, R. *et al.* MacaquePose: A Novel “In the Wild” Macaque Monkey Pose Dataset for Markerless Motion  
551 Capture. *Front. Behav. Neurosci.* **14**, 268 (2021).
- 552 36. Graving, J. M. *et al.* DeepPoseKit, a software toolkit for fast and robust animal pose estimation using deep learning.  
553 *eLife* **8**, e47994 (2019).
- 554 37. Cao, J. *et al.* Cross-Domain Adaptation for Animal Pose Estimation. *Proc. IEEE Int. Conf. Comput. Vis.* **2019-**  
555 **October**, 9497–9506 (2019).
- 556 38. Yu, H. *et al.* AP-10K: A Benchmark for Animal Pose Estimation in the Wild. Preprint at  
557 <https://doi.org/10.48550/ARXIV.2108.12617> (2021).
- 558 39. Liu, H., Liu, F., Fan, X. & Huang, D. Polarized self-attention: Towards high-quality pixel-wise mapping.  
559 *Neurocomputing* **506**, 158–167 (2022).
- 560 40. Xu, Y., Zhang, J., Zhang, Q. & Tao, D. ViTPose: Simple Vision Transformer Baselines for Human Pose  
561 Estimation. Preprint at <https://doi.org/10.48550/arXiv.2204.12484> (2022).
- 562 41. Yang, S., Quan, Z., Nie, M. & Yang, W. TransPose: Keypoint Localization via Transformer. Preprint at  
563 <https://doi.org/10.48550/ARXIV.2012.14214> (2020).
- 564 42. Korschens, M. & Denzler, J. ELPephants: A Fine-Grained Dataset for Elephant Re-Identification. in *2019*  
565 *IEEE/CVF International Conference on Computer Vision Workshop (ICCVW)* 263–270 (IEEE, Seoul, Korea  
566 (South), 2019). doi:10.1109/ICCVW.2019.00035.
- 567 43. Xu, Y., Zhang, J., Zhang, Q. & Tao, D. ViTPose++: Vision Transformer for Generic Body Pose Estimation.  
568 Preprint at <https://doi.org/10.48550/ARXIV.2212.04246> (2022).
- 569 44. Price, J. (July 25, 2024). The kangaroo: All you need to know about Australia's most iconic animal - and its famous  
570 hop [Image]. *DiscoverWildlife*. <https://www.discoverwildlife.com/animal-facts/mammals/kangaroo-facts>.
- 571 45. Yu, H. *et al.* AP-10K: A Benchmark for Animal Pose Estimation in the Wild. (2021).
- 572 46. Ferreira, B. (2020). African Wildlife, version 1. Retrieved January 3, 2025 from  
573 <https://www.kaggle.com/datasets/biancaferreira/african-wildlife>.
- 574 47. Banerjee, S. (2022). Animal Image Dataset (90 Different Animals), version 5. Retrieved January 3, 2025 from  
575 <https://www.kaggle.com/datasets/iamsouravbanerjee/animal-image-dataset-90-different-animals>
- 576 48. Trivedi, Y. V. (2021). animals\_5, version 1. Retrieved January 3, 2025 from  
577 <https://www.kaggle.com/datasets/yttrivedi1/animals-5>.
- 578 49. Alessio, C. (2020). Animals-10, version 1. Retrieved January 3, 2025 from  
579 <https://www.kaggle.com/datasets/alessiocorrad099/animals10>.
- 580 50. Jamil, S. (2021). Endangered Animals, version 1. Retrieved January 3, 2025 from  
581 <https://www.kaggle.com/datasets/sonain/endangered-animals>.
- 582 51. Antoreepjana. (2021). IUCN Animals Dataset, version 1. Retrieved January 3, 2025 from  
583 <https://www.kaggle.com/datasets/antoreepjana/iucn-animals-dataset>.
- 584 52. Sahovic, E. (2021). wild\_cats, version 1. Retrieved January 3, 2025 from  
585 <https://www.kaggle.com/datasets/enisahovi/cats-projekat-4>.
- 586 53. Vivek. (2022). Asian vs African Elephants, version 1. Retrieved January 3, 2025 from  
587 <https://www.kaggle.com/datasets/vivmankar/asian-vs-african-elephant-image-classification/data>.
- 588 54. Bengio, Y., Louradour, J., Collobert, R. & Weston, J. Curriculum learning. *ACM Int. Conf. Proceeding Ser.* **382**,  
589 (2009).
- 590 55. Paszke, A. *et al.* PyTorch: An Imperative Style, High-Performance Deep Learning Library. Preprint at  
591 <https://doi.org/10.48550/ARXIV.1912.01703> (2019).
- 592 56. Kingma, D. P. & Ba, J. Adam: A Method for Stochastic Optimization. Preprint at  
593 <https://doi.org/10.48550/ARXIV.1412.6980> (2014).
- 594 57. Kirkpatrick, J. *et al.* Overcoming catastrophic forgetting in neural networks. *Proc. Natl. Acad. Sci.* **114**, 3521–3526  
595 (2017).
- 596 58. Ronchi, M. R. & Perona, P. Benchmarking and Error Diagnosis in Multi-Instance Pose Estimation. Preprint at  
597 <https://doi.org/10.48550/ARXIV.1707.05388> (2017).
- 598 59. COCO - Common Objects in Context. Available at <https://cocodataset.org/#keypoints-eval>. (Accessed: 20-Dec-2021).
- 599